

NEURAL LANGUAGE TRANSLATION

LAUREN O'DONNELL

DS 6050 DEEP LEARNING

FQ JIXDYMEBSLJBWXDUNL
GFBVWLC TFP OIZQAYWHAT
MYVLOYFJR CVUNIJ PNJHI
WZUXQURAXIOMVMVOFTDC
VYCDYCJKMOPXEFRS PCOB
KBJIMVKIVAGVGRQNT EZH
ZHYBSECNIMDGOMFVETO E
CIPUYKFIXOCTFZCHJEAR
YKRVEGICRLXCLKLCTRD
QLGZRWF PFOEIYFVRMZHX
RPZYDUIVTEAXLJWSIRUC
JLAVMPLOTYCKIBQYWYPK
BPFRDJTVAQIFSTZVFMJC
SYECVINGFBRNYUCBSNTD
CFIBRMSZJEDXRWTKADFE



AGENDA

- Introduction
- Data
- Methods
- Results
- Conclusions



INTRODUCTION

- Neural Machine Translation (NMT) is a branch of Machine Translation (MT) which falls under Natural Language Processing (NLP)
 - Aims to translate natural language sentences
- Used in Google Translate, Amazon Translate, and many other common translation services & applications
- Many use cases: book translations, caption translation, research paper/website translations, etc.
- Consistency is an outstanding challenge!



INTRODUCTION: A BRIEF HISTORY

- Early approaches relied on hand-crafted translation rules and linguistic knowledge
 - Challenges: language is COMPLEX and there are many language IRREGULARITIES
- Statistical Machine Translation (SMT)
 - Learns latent structures
 - Uses discrete symbolic representations
 - Separately tuned components
 - Challenges: unable to model long-distance dependencies between words; poor quality
- Neural Machine Translation (NMT)
 - Single large neural network to model entire translation process
 - Uses continuous symbolic representations
 - End-to-end training
 - Probabilistic framework



INTRODUCTION: NMT ARCHITECTURE

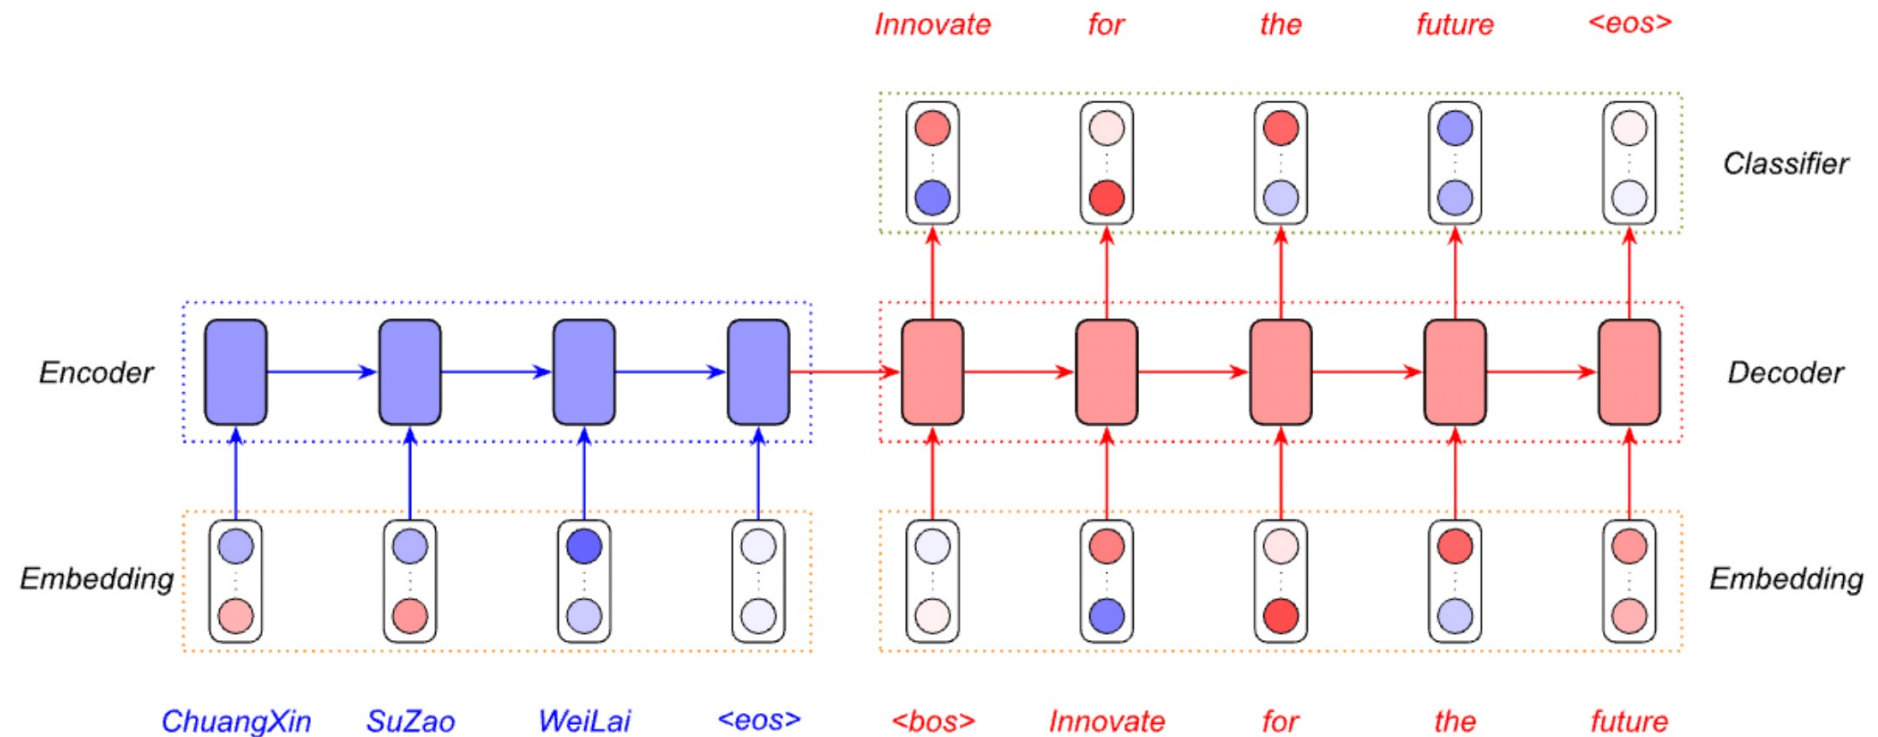
- NMT is a sequence-to-sequence model
 - Given a source sentence $x = \{x_1, \dots, x_s\}$ and a target sentence $y = \{y_1, \dots, y_T\}$, the chain rule and the conditional distribution can factorize from left-to-right (L2R) as:

$$P(y|x) = \prod_{t=1}^T P(y_t | y_0, \dots, y_{t-1}, x)$$

- Three main sections:
 - Embedding layers
 - Encoder and decoder networks
 - Classification layer
- Common encoders and decoders include RNN, CNN, and SAN



INTRODUCTION: NMT ARCHITECTURE



INTRODUCTION:

ENCODERS & DECODERS

RECURRENT NEURAL NETWORKS

- Pros
 - Extremely powerful
 - Infinite receptive fields
 - Does not require additional padding or masking due to sequential nature
- Cons
 - Not compatible with GPU/TPU
 - Severe vanishing and exploding gradient problems

CONVOLUTIONAL NEURAL NETWORKS

- Pros
 - Compatible with GPU/TPU
- Cons
 - Requires additional padding and masking to prevent the network from seeing future words
 - Limited receptive field

SELF ATTENTION NETWORKS

- Pros
 - Compatible with GPU/TPU
 - Infinite receptive fields
- Cons
 - Requires additional padding and masking to prevent the network from seeing future words



INTRODUCTION: GOALS

- Explore and compare two techniques and architectures for NMT:
 - RNN and SAN



DATA

- English-Portuguese datasets
- Contains 168,903 total records
 - Mix of words (like exclamations) and sentences
 - 80-20 training to validation set split
- Features are the words!
 - Varies sentence to sentence and phrase to phrase
- [Kaggle Notebook](#) pre-processed the data and generated an RNN model utilizing GRU and Attention Mechanism



METHODS:

FUNDAMENTAL ARCHITECTURES

RNN: GATED RECURRENT UNIT (GRU) WITH ATTENTION MECHANISM

- RNN model
- Basic encoder-decoder
- Consists of 1 encoder GRU layer and 1 GRU decoder layer with an attention mechanism


SAN: BIDIRECTIONAL ENCODER REPRESENTATIONS FROM TRANSFORMERS (BERT)

- SAN model
- Additional sinusoid-style position encoding added to input embedding
- Consists of 6 encoder layers and 6 decoder layers



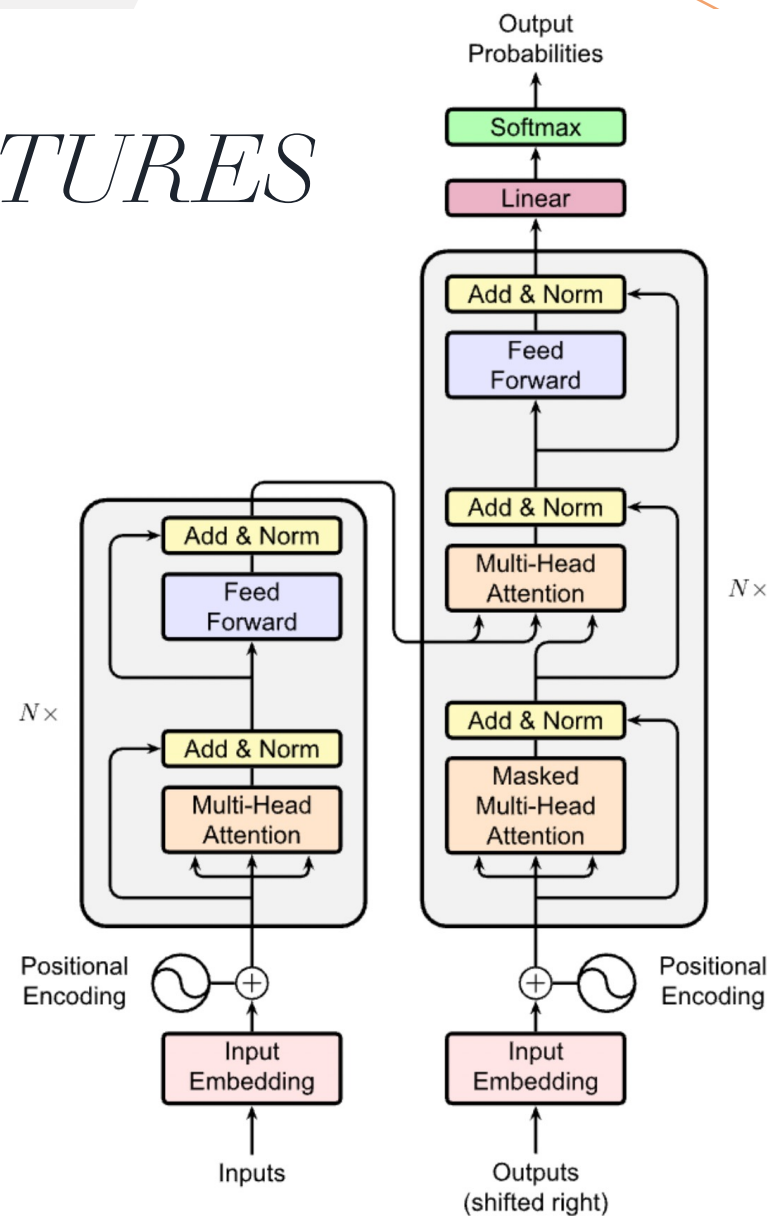
METHODS: *FUNDAMENTAL ARCHITECTURES*

RNN: GATED RECURRENT UNIT (GRU) WITH ATTENTION MECHANISM



METHODS: *FUNDAMENTAL ARCHITECTURES*

TRANSFORMERS



METHODS: HYPERPARAMETERS

- Learning Rate: 0.000001
- EPOCHS: 10
- Batch Size: 100
- Activation Function: softmax for RNN, sigmoid for SAN
- Dropout Rate: 0.1

```
class Encoder(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, enc_units, batch_sz):
        super(Encoder, self).__init__()
        self.batch_sz = batch_sz
        self.enc_units = enc_units
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = gru(self.enc_units)

    def call(self, x, hidden):
        x = self.embedding(x)
        output, state = self.gru(x, initial_state = hidden)
        return output, state


    def initialize_hidden_state(self):
        return tf.zeros((self.batch_sz, self.enc_units))
```

```
class BERT_Encoder(tf.keras.Model):
    def __init__(self, max_length, num_classes):
        super(BERT_Encoder, self).__init__()
        self.max_length = max_length
        self.num_classes = num_classes
        self.bert = TFBertModel.from_pretrained('bert-base-uncased')
        self.dropout = tf.keras.layers.Dropout(0.1)
        self.classifier = tf.keras.layers.Dense(num_classes, activation='sigmoid')

    def call(self, inputs, **kwargs):
        _, pooler_output = self.bert(inputs, **kwargs)
        x = self.dropout(pooler_output, training=kwargs.get('training', False))
        x = self.classifier(x)
        return x
```






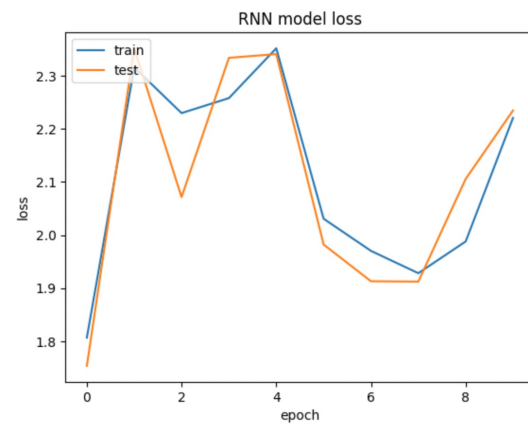
Input: tom still does that quite often
Predicted translation: Tom ainda faz isso com muita frequência
Actual translation: Tom ainda faz isso com bastante frequência



RESULTS

- RNN model:
 - Training accuracy: 78.48%
 - Validation accuracy: 78.20%
- BERT model:
 - Training accuracy: 14.86%
 - Validation accuracy: 14.42%





RESULTS



CONCLUSIONS

RESEARCH CONCLUSIONS

- RNN model performed better
 - Training accuracy: 78.48%
 - Validation accuracy: 78.20%
- Language translation is COMPLEX
 - Easy to overfit

FUTURE RESEARCH & RECOMMENDATIONS

- Significantly more compute resources are required
- Requires more training time
- How do other languages either more similar or different affect results? Ex.:
 - Spanish to Portuguese
 - English to Arabic



REFERENCES

- Zhixing Tan, Shuo Wang, Zonghan Yang, Gang Chen, Xuancheng Huang, Maosong Sun, and Yang Liu. Neural machine translation: A review of methods, resources, and tools. *AI Open*, 1:5–21, 2020.
- Kalchbrenner, Nal and Espeholt, Lasse and Simonyan, Karen and Oord, Aaron van den and Graves, Alex and Kavukcuoglu, Koray. Neural machine translation in linear time. *arXiv preprint arXiv:1610.10099*, 2016.
- *Manning Luong, Pham. Effective approaches to attention-based neural machine translation. <https://paperswithcode.com/paper/effective-approaches-to-attention-based>, 2015.*
- A history of machine translation from the Cold War to deep learning. Freecodecamp, <https://www.freecodecamp.org/news/a-history-of-machine-translation-from-the-cold-war-to-deep-learning-f1d335ce8b5/#:~:text=On%20January%207th%201954%2C%20at,the%20first%20time%20in%20history>, 2018.
- Rani Horev. BERT Explained: State of the art language model for NLP. Towards data science, <https://towardsdatascience.com/bert-explained-state-of-the-art-language-model-for-nlp-f8b21a9b6270>, 2018.
- <https://huggingface.co/>
- <https://www.kaggle.com/code/walterfly/use-hugging-face-s-tensorflow-2-transformer-models>
- https://www.tensorflow.org/text/tutorials/nmt_with_attention
- <https://www.tensorflow.org/text/tutorials/transformer>
- <https://www.kaggle.com/code/nageshsingh/neural-machine-translation-attention-mechanism/notebook>